

Assignment - 1

1. Modern systems still rely on OS because:
 - Abstraction: OS hides hardware complexity (device, CPU, memory) so apps run on many platforms.
 - Resource management: OS schedules CPU, memory, I/O.
2. Real time embedded OS (RTOS) is used because of predictable timing, low power footprint, small memory / CPU, deterministic interrupt handling.
3. While building a new kernel we should avoid microkernels if raw performance is the only goal. Microkernels give modularity and safety but add IPC overhead because many services run in user space.
4. OS structure doesn't matter as long as process runs. Refute because of structure affects:-
 - Performance
 - Reliability / security
 - Maintainability and extensibility
 - Feature support.
- 5.i). PCB stores saved registers, PC, stack pointer, state bits
 - PCB's saved PC points to an unexpected address, the next resume will jump to wrong code.
- ii). Save current CPU context into its PCB
 - Change process state to waiting and enqueue on wait queue.
 - Update memory mapping
 - Scheduler selects next ~~memory~~ process

- iii). Use asynchronous non-blocking I/O when you don't want the task to block and you want responsiveness.
- If task must wait for I/O result before proceeding, use blocking synchronous ~~set~~ calls.

6. Save state = 2ms

Load state = 3ms

Scheduler overhead = 1ms

a) Total context switching time = $2 + 3 + 1 = 6\text{ms}$

b) Impact on multitasking:

→ Every switch costs 6ms of CPU. High context-switch rate reduces effective CPU for user work.

→ Frequent switching increases cache.

7. Single threaded total time = 40s

Threads per process = 2

• Estimate execution time = $\frac{40}{2} = 20\text{ seconds}$

• Multithreading helps allows parallelism on multiple core and hides latency by overlapping I/O with computation.

8. Process P₁ P₂ P₃ P₄
Time(ms) 5 3 8 6

a) Gantt chart

• FCF S

P ₁	P ₂	P ₃	P ₄
0	5	8	16

• SJF

P ₂	P ₁	P ₄	P ₃
0	3	8	14

- Round Robin (Quantum = 4ms)

P ₁	P ₂	P ₃	P ₄	P ₁	P ₃	P ₄
0	4	7	11	15	16	20 22

b) Waiting time = Start time - arrival time
 $TAT = \frac{\text{burst}}{\text{quantum}}$

Turnaround time = Finish time - arrival time

FCFS : P₁ = 5, P₂ = 8, P₃ = 16, P₄ = 22

$$\begin{aligned} TAT &= 5 + 8 + 16 + 22 = 51 \\ &= \frac{51}{4} = 12.75 \end{aligned}$$

Waiting time = TAT - burst
 $= 0 + 5 + 8 + 16 = 29.$

$$\text{Avg. WT} = \frac{29}{4} = 7.25$$

SJF P₂ = 3, P₁ = 8, P₄ = 14, P₃ = 22

$$\begin{aligned} TAT &= 8 + 3 + 22 + 14 = 47 \\ &= \frac{47}{4} = 11.75 \end{aligned}$$

Waiting time = TAT - burst

$$\begin{aligned} &= 8 - 5 = 3 \\ &= 3 - 3 = 0 \\ &= 22 - 8 = 14 \\ &= 14 - 6 = 8 \end{aligned} \quad \left. \begin{array}{l} \\ \\ \end{array} \right\} \text{Sum} = 25$$

$$\text{Avg. WT} = \frac{25}{4} = 6.25$$

Round Robin P₁ = 16, P₂ = 7, P₃ = 20, P₄ = 22

$$TAT = 16 + 7 + 20 + 22 = 65$$

$$\text{Avg. TAT} = \frac{65}{4} = 16.25$$

Waiting time > TAT - burst

$$\begin{aligned} 16-5 &= 11 \\ 7-3 &= 4 \\ 20-8 &= 12 \\ 22-6 &= 16 \end{aligned} \quad \left. \begin{array}{l} \text{sum} = 43 \\ \text{TAT} \end{array} \right\}$$

$$\text{Avg. WT} = \frac{43}{4} = 10.75$$

c) SJF gives lowest average turnaround 11.75 and lowest avg. waiting 6.25 among the three.

9.1(a) Microkernel or modular layered architecture is preferred for scalability and security.

- Microkernel : small trusted core, most services run in user space.
- Layered modular : It is designed also aids maintainability and clear separation.

b). Isolation : Each virtual machine has separate guest OS and virtual hardware.

• Management : Virtual machine support snapshots, cloning, live migration, templates for rapid deployment.

• Resource optimisation : Hypervisors multiplex CPU/memory/disk across virtual machine support overcommit, dynamic scaling and scheduler-level QoS.

- i) Ensuring high priority tasks are handled without delay.
- Priority-based scheduling : Assign high priority to intrusion detection tasks.
 - Real-time scheduling : Use ~~the~~ Real time scheduler for critical tasks so deadlines are met.
 - Priority inheritance : Avoid priority inversion when lower-priority tasks hold resources needed by high-priority tasks.
- b). Rate Monotonic Scheduling (RMS) : Good for periodic hard-real time tasks with fixed periods.
- Earliest Deadline First (EDF) : Optimal for dynamic deadlines and mixed workloads - better CPU utilization if tasks have different deadlines.
 - Hybrid approach : Use real-time scheduler (RMS) for safety/critical tasks.

Saturday
20/11/25