

Assignment - 4

1. Two clerks share a single physical ledger book. Both read the current balance, each adds a deposit amount on their copy, and both write back. If one writes after the other, one deposit is lost - result depends on timing (a race). Mutual exclusion fix: Use a checkout for the ledger. A clerk must acquire the lock before reading/updating & release it after.

Q.5

Part B

$$S_1 : P_1 \rightarrow P_2, P_3 \rightarrow P_4$$

$$S_2 : P_2 \rightarrow P_5, P_5 \rightarrow P_6$$

$$S_3 : P_6 \rightarrow P_1$$

$$\begin{array}{ll} \text{a) Edges} = & \\ P_1 \rightarrow P_2 & P_5 \rightarrow P_6 \\ P_3 \rightarrow P_4 & P_6 \rightarrow P_1 \\ P_2 \rightarrow P_5 & \end{array}$$

- b) Deadlock detection: look for cycles: $P_1 \rightarrow P_2 \rightarrow P_5 \rightarrow P_6 \rightarrow P_1$ forms a cycle. Process involved: P_1, P_2, P_5, P_6 ($P_3 \rightarrow P_4$ is separate & not in cycle.)

- c) Suggested distributed algorithm: Chandy - Misra - Haas (edge chasing).

$$7.a) \text{Expected file access time calculation: } \text{Prob(local)} = 1 - 0.3 \\ = 0.7$$

$$\begin{aligned} \text{Expected time} &= 0.7 \times 5 \text{ ms} + 0.3 \times 25 \text{ ms} \\ &= 3.5 \text{ ms} + 7.5 \text{ ms} = 11.0 \text{ ms} \end{aligned}$$

- b) Caching strategy recommendation because it is frequently

accessed remote files become local after first access, reducing remote fetches and average latency

- 8.a) Proposed mix (10 s second period): Full checkpoint once every 10s (1 full checkpoint per 10s)
Cost = 200ms

Incremental checkpoints: taken every 1s b/w fulls
(at t=1, 2, ..., 9) \rightarrow 9 incrementals \times 50ms = 450ms

Total overhead in 10s = 200 + 450 = 650ms
i.e. = 0.65s of CPU time per 10s period.)

- b) Reasoning: RPO = 1s requires that at most 1 second of new state is lost on failure, so taking incremental checkpoints every 1s satisfies RPO.

- 9.a) Extremely bursty request patterns causing hotspots.
Need for low latency (geo-locality) while balancing global inventory/consistency. Heterogeneous servers, network latency, and autoscaling delays. Suggested algorithm: 'Power of Two Choices' (randomized load balancing).

Set 20/11/22

- b) Active-active multi-region deployment with geo-replication
critical state replicated synchronously (or semi-synchronously)
across at least two regions for services requiring low RPO; less-critical data can use asynchronous replication. Use distributed consensus (Raft / Paxos) for metadata & leader election to enable fast failover.

Part - A

1. Race condition : Two people try to write different notes on the same sticky note at the same time. If both read current text, append their own & write back without coordination, one person's update can overwrite other's - result depends on timing
2. Implementation complexity : Peterson's solution is simple algorithmically ~~for~~ but limited to two processes and is tricky to extend. Semaphores provides a general high level abstraction suitable for many process.
~~Hardware dependency : Peterson's relies on sequential consistency and atomic writes to shared variables - it may fail on modern weak memory multiprocessors without memory fences~~
3. Monitors provide structured mutual exclusion plus condition variables built in. In a multi-core system, monitors can use efficient per-monitor locks & condition queues, reducing context switch churn & avoiding busy-waiting.
4. Writers starve if readers continuously arrive because now readers can keep entering and block waiting writers forever. We can prevent it by using a fair policy so that once a writer arrives it eventually gets priority and waits readers are blocked until writer completes.
5. Eliminating Hold and Wait typically forces process to acquire all required resources at once ^{or} for release held resource before requesting new ones). Practical drawback : poor resource utilization and reduced concurrency process may